

# Supporter as Studio: Distributed AI Media Production in the 2026 Los Angeles Mayoral Race

Lance Rossi

## Abstract

This paper uses the 2026 Los Angeles mayoral campaign of Spencer Pratt to examine an emerging configuration of political media production: cinematic, culturally targeted campaign video generated by supporters using consumer-accessible generative AI tools and amplified through the candidate's own social channels. Drawing on Chadwick's hybrid media system framework, extended with an explicit production layer, and on Penney's account of the citizen marketer, the paper offers a two-tier production typology coupled to an amplification stack, anchored in four artifacts. It argues that supporter-generated AI media partially offsets resource asymmetries in the attention economy while leaving electoral outcomes untouched "Pratt achieved attention parity without outcome parity, finishing third behind a conventional candidate who matched his fundraising without any AI ecosystem" and identifies two characteristic failure modes: cultural-fit misfit and political-economy misfit. The June 2 primary result, including Pratt's fourteen-year high-water mark for a GOP-aligned candidate, motivates a refinement of the persuasion-versus-consolidation question into a persuasion/consolidation/mobilization triad for the 2028 research program.

## Research problem

In the spring of 2026, the Los Angeles mayoral campaign of Spencer Pratt "reality television personality, registered Republican in a nominally nonpartisan race, and challenger to incumbent Karen Bass" became the most visible recent example of a configuration of political communication that the field is only beginning to take seriously: a campaign whose most cinematically ambitious media is produced not by the campaign but by ideologically aligned supporters using consumer-accessible generative AI tools, and is circulated principally through the candidate's amplification of supporter-generated content. This paper asks: how should the field of political communication theorize the production-side change that generative AI introduces into the hybrid media system, and what does the Pratt case suggest about the durability, generalizability, and characteristic vulnerabilities of the configuration?

This is a single-case, theory-building study, and it claims no more than that. The Pratt case is unusually clean for the purpose "the supplier population is small enough to name, the artifacts few enough to describe individually, the press response documented across at least five distinct media frames. Whether comparable ecosystems are emerging around

other 2026 campaigns is plausible given the technology's accessibility, but it is an empirical question this paper does not answer; the case selection logic is that of the revelatory case, not the representative one. What a single clean case can do is generate concepts precise enough to be tested against the larger population of cases the 2028 cycle will supply.

The paper offers four contributions: (1) a typology distinguishing two tiers of media *production* coupled to an *amplification stack*, resisting the campaign-versus-supporter binary common in current commentary; (2) an extension of Chadwick's (2017) hybrid media system framework with an explicit production layer; (3) two distinct failure modes "cultural-fit misfit and political-economy misfit" that constrain the configuration; and (4) a post-primary update that converts the paper's central hedge about virality and votes into a finding, introducing *mobilization* as a category between persuasion and consolidation.

## Theoretical framing

Four literatures are relevant. Chadwick's (2017) *hybrid media system* is the master frame: contemporary political communication operates through the interaction and mutual remediation of older and newer media logics, neither of which simply supersedes the other. Bennett and Segerberg's (2013) *connective action* characterizes the loose, identity-anchored coordination of personalized political expression at scale, and offers part of the account of why supporters produce political media without coordination. The early empirical literature on generative AI in campaigns "notably the Center for Media Engagement's "Political Machines" report (2024), the Stanford / U. Chicago Harris working group recommendations (Persily, Tucker, et al., 2024), and Farooq, van den Hoogen, Tulin, and de Vreese's (2025) *International Journal of Press/Politics* study of European citizens' perceived versus actual exposure to AI election content" frames the technical and audience-side change.

The fourth literature is the one this configuration is most easily mistaken for superseding: the scholarship on participatory political culture and citizen marketing. Jenkins's (2006) account of convergence culture established that audiences appropriate, remix, and redistribute media texts as acts of identity and affiliation; Penney's *The Citizen Marketer* (2017) theorized, well before consumer generative video existed, the supporter who voluntarily produces and circulates persuasive political media as an extension of selfhood rather than as an agent of the campaign. On Penney's account, the Curran and Alvarez artifacts described below are not a new species of political behavior "they are citizen marketing. What is genuinely new is not that supporters produce; it is *what supporters can now produce*. The 2017 citizen marketer made memes, hashtags, and selfie videos; the 2026 citizen marketer makes cinematic, broadcast-quality narrative video at effectively zero marginal cost. The change is in production quality and grammar, not in the existence of the practice, and the paper's novelty claim is calibrated accordingly. A second calibration follows from the celebrity-politics literature: Pratt arrived with fifteen years of accumulated parasocial audience, reality-television media fluency, and entertainment-

industry donor networks. The density of his supporter-production ecosystem may therefore be partly an artifact of fan culture around a celebrity candidate; a non-celebrity candidate with identical tools might observe a much thinner Tier II. This confound cannot be resolved within a single case, and it is flagged here rather than resolved.

The synthesis this paper proposes is that Chadwick's framework, while still the right master frame, needs an explicit *production layer*. In the original formulation, who can produce broadcast-quality political video is largely treated as exogenous: campaigns produce, journalists produce, citizens produce, and the system distributes. What the Pratt case suggests is that the production end is no longer exogenous. Generative video tools accessible to consumers in 2026 (Charlie Curran's reporting indicates use of Seedance 2.0, a tool that was not consumer-accessible at the time of the 2024 US election) collapse the cost of producing cinematic political imagery by an order of magnitude, opening the production layer to a much larger supplier population. This is an infrastructural change distinct from changes in distribution or audience composition, and it requires its own theoretical treatment.

The connective action frame extends naturally to the new condition, with one qualification: in a 2013 connective action episode, production cost was roughly even across participants (everyone tweets, everyone memes). In a 2026 connective action episode involving generative AI, production cost is technically low across the board but production *skill* is unevenly distributed. This means that supporter-generated AI content is internally stratified by the supporter's existing production capacity and cultural knowledge — a stratification that the Pratt case makes visible.

## The Pratt media ecosystem

Spencer Pratt announced his candidacy for LA mayor on January 7, 2026 — the first anniversary of the Palisades fire that destroyed his family's home. Polling traced an atypical trajectory: in March, Bass led with 20%, Pratt at 10%, Raman at 9%, undecideds at 51%; by mid-May, Bass was at 30%, Pratt at 22%, Raman at 19%, undecideds at 16% (KTLA / Emerson, May 2026). The fundraising picture carries theoretical weight later in the paper. Pratt raised approximately \$530,000 following his February 2026 committee filing — closing the cycle-on-cycle gap with the incumbent (Bass raised \$494,734 in the current cycle) while Bass retained the institutional money advantage (\$2.8M total since 2024, \$2.3M cash on hand) (TheWrap, May 2026). Critically, Councilwoman Nithya Raman *also* raised approximately \$530,000 in the same period — matching Pratt dollar for dollar without any supporter-generated AI ecosystem, a fact the takeaways section confronts directly. Pratt was widely assessed as the strongest performer in the mid-May NBC4/Telemundo 52 debate (*Washington Post* Ripple; *The Hill* opinion, May 2026); Bass subsequently withdrew from a scheduled second debate, traveling to Sacramento for state funding meetings instead.

Four media artifacts anchor the analysis.

The *Batman / Dark Knight Rises* ad (early May 2026) was produced by Charlie Curran of Menace Studios, an LA-based AI filmmaking shop. The ad depicts LA as Gotham under aristocratic misrule — Bass as Joker, Newsom as cake-eating Marie Antoinette, Harris vodka-chugging — with Pratt as Batman leading citizens to throw tomatoes at the elites, echoing the mob-rule trial scenes of *The Dark Knight Rises*. Within roughly twenty-four hours it had drawn 3.9 million views on X and 10,000 reposts (Hollywood Reporter, May 2026), amplified by Jeb Bush (“maybe the best political ad of the year”), Ted Cruz, and Matt Gaetz. Pratt did not commission it; Curran has no formal campaign affiliation; Pratt reposted it.

The “*Spencer, Saca La Bassura*” video (released March 30, 2026) was produced by Adrian E. Alvarez, a Cuban-American account holder who runs Latinos Por Pratt and splits his time between Miami and Los Angeles. The AI-animated salsa music video puns on Bass’s surname and the Spanish word for trash; Pratt wheels out a trash bin containing Bass. Gustavo Arellano of the *Los Angeles Times* argued in April 2026 that the video’s salsa instrumentation, Cuban Spanish lexicon, and Miami sensibility misfire in a city whose Latino population is overwhelmingly Mexican — a corrido in Mexican Spanish would have fit; a salsa in Cuban Spanish does not. The artifact is, on the Arellano critique, the cleanest available evidence of *cultural-fit failure* in supporter-generated AI media.

The “*Fresh Prince of Bel-Air*” parody (May 18, 2026) was posted to Pratt’s own Instagram and is campaign-produced content. Pratt raps a parody of the *Fresh Prince* theme telling the story of the Palisades fire displacing him, ending the video pulling up to the Hotel Bel-Air — an ending that triggered a backlash and a KTLA story in which Pratt denied living at a luxury hotel.

The “*They Not Like Us*” video (April 2026) was commissioned by Pratt’s campaign (per NBC News reporting). It plays on Kendrick Lamar’s hit, films Pratt outside Bass’s mansion and Raman’s “\$3 million mansion,” and pivots to Pratt’s Airstream on the scorched lot where his Palisades home stood.

These four artifacts — two supporter-generated, two campaign-produced (one candidate-made, one campaign-commissioned) — span the typology this paper develops. They traveled through a recognizable amplification stack: supporter to candidate to ideologically aligned elites (Bush, Cruz, Gaetz, Joe Rogan, Donald Trump) to entertainment trade press (Hollywood Reporter, TMZ, Parade) to mainstream broadcast and digital news (CNN, NBC, NPR, ABC7, KTLA, *Los Angeles Times*) to opposition response (Bass on CNN describing the artifacts as “a very dangerous trend”). Each layer compresses the time between artifact and event; each layer also reshapes the artifact’s meaning, as the entertainment press encodes spectacle and the political press encodes a “new political reality.”

## Two production tiers and an amplification stack

An earlier version of this analysis described a “three-tier typology” of media production. That formulation contained a category error worth correcting publicly, because the

correction is instructive. Tiers I and II classified artifacts by *production provenance* “who made the thing. The former Tier III classified actors by *circulation behavior* “who moved the thing. Folding a distribution category into a production typology conflates two ontologically distinct questions, and a typology that cannot keep its own dimensions apart will not survive contact with a larger case population.

The revised architecture is therefore a two-tier production typology coupled to an amplification stack. *Tier I* is campaign-produced content “material produced or directly commissioned by the candidate or formal campaign (the Fresh Prince parody, the “They Not Like Us” video, the official channels). *Tier II* is supporter-generated AI content “material produced by ideologically aligned supporters without formal campaign coordination (Curran’s Batman ad and his lightsaber-duel and John Wick “style ads; Alvarez’s Saca La Bassura video and the Wanted Out of Office posters). The *amplification stack* is not a tier at all: it is the ordered sequence of circulation roles “candidate, aligned elites, entertainment trade press, mainstream political press, opposition response “through which artifacts from either tier travel.

The architectural separation does real analytic work. Production provenance is a property of the *artifact*, fixed at the moment of creation: the Batman ad is Tier II forever, regardless of who circulates it. Position in the amplification stack is a property of an *actor’s role in a particular circulation episode*, and the same actor can occupy different positions for different artifacts “Pratt is a Tier I producer of the Fresh Prince parody and a first-rung amplifier of the Batman ad. What previously looked like the typology “leaking” “Pratt reposting Tier II content from Tier I channels, Cruz becoming a secondary distributor “is not leakage at all once the dimensions are kept distinct; it is the expected behavior of a system in which any actor may move any artifact. Boundary-crossing between categories would be a defect of the typology; movement of artifacts across circulation roles is the phenomenon the typology exists to describe. The separation is also faithful to Chadwick’s own logic: the hybrid media system is precisely a model in which production logics and circulation logics interpenetrate without becoming identical, and the configuration analyzed here is a system of that kind, not a sequence. The two tiers and the stack are functionally interdependent “the Batman ad would not have reached 3.9 million views without elite and trade-press amplification, and that amplification would not have occurred without the supporter-generated artifact to amplify.

## Key takeaways

*Production-cost asymmetry buys attention parity, not outcome parity.* The strongest version of the resource-asymmetry claim “that supporter-generated AI media offsets the incumbent’s money advantage “is contradicted by the case’s own comparison candidate, and the paper is better for confronting it. Raman raised the same \$530,000 as Pratt with no AI supporter ecosystem, no viral artifacts, and a debate performance the *Los Angeles Times* scored as weak “and Raman made the runoff while Pratt did not. What supporter-generated AI media demonstrably bought Pratt was *media reach*: national coverage across five press frames, elite amplification, and a polling trajectory no third-

place local candidate ordinarily enjoys. What it did not buy was the outcome a conventional candidate achieved with equivalent money and none of the apparatus. The refined claim is narrower and more defensible: the production-layer change partially offsets resource asymmetry *in the attention economy* “the air war” while leaving ground-game, institutional-endorsement, and coalition advantages untouched. This is a refinement, not a retreat; a theory that predicted attention effects and disclaimed outcome effects is exactly what the evidence now supports.

*Production-skill stratification is the configuration’s first characteristic vulnerability.* Inside Tier II, the cost of *production* approaches zero, but the cost of *production skill* and *cultural knowledge* remains. The Curran/Alvarez contrast is the cleanest available illustration: Curran is a commercial filmmaking veteran producing for a community he belongs to; Alvarez is, on available reporting, working alone and producing for a community he does not belong to. The Arellano critique identifies the failure mode likely to scale in proportion to supporter participation: cultural-fit failure when culturally distant supporters produce content for communities they don’t understand.

*Political-economy misfit is a second, distinct failure mode.* Raman’s campaign articulated a critique that is not about cultural fit at all: Pratt, she charged, promotes AI-generated media in the one American city whose film and television workforce is being displaced by precisely that technology “and whose largest crafts union, IATSE, was funding his opponent. This is a misfit between the *means of the message* and the *material interests of the audience*, analytically distinct from the Saca La Bassura problem: a supporter-generated artifact can be culturally pitch-perfect and still activate this critique, because the objection attaches to the production technology itself. The critique is also notable as a left-coded counter-frame. Most contemporaneous criticism of the Pratt ecosystem ran through disinformation and “AI slop” frames; the labor frame is the one most likely to travel in Democratic-coalition cities, and the one campaigns adopting the configuration will most need an answer to.

*Attribution and accountability are unresolved “and the regulatory question is live.* When Bass on CNN described the AI artifacts as “a very dangerous trend” (citing a clip showing her and Newsom drowning in a reservoir), she raised an attribution question the field has not yet answered. Pratt’s description of the videos as “fan-made” (in a since-deleted Instagram clip, per NPR) is precisely the attribution-evading move the configuration enables. The supporter is not inauthentic; the content is not strictly campaign content; the campaign has plausible deniability. And yet the candidate’s amplification gives the content a status it would not otherwise have, and the candidate clearly benefits from its reach. The configuration creates, structurally, a space of *attributable-but-deniable* campaign content. That space has a legal dimension this paper can only flag. Supporter-generated content amplified by the candidate sits in contested territory around coordination doctrine and in-kind contribution rules: if a candidate’s repost constitutes adoption, the value of the creative labor behind the artifact is arguably reportable; if not, the configuration offers campaigns a channel for professional-quality media wholly outside disclosure regimes. Neither California’s Fair Political Practices Commission nor the Los Angeles City Ethics



Commission appears, as of this writing, to have issued guidance on candidate amplification of supporter-generated AI media, and the valuation question “what is the in-kind value of a video that cost almost nothing to render but draws on substantial skill?” has no settled answer. The regulatory vacuum is itself a finding: the configuration scales faster than the rules that would govern it.

*The virality-to-votes link remains weak “by design of the analysis, and now by evidence.* Pratt’s polling movement from March (10%) to May (22%) coincided with the undecided pool collapsing from 51% to 16%, debate-performance gains, and the cumulative effect of his loss-of-home narrative; it was never possible to attribute the share gain principally to supporter-generated AI media. The paper makes claims about production-layer change and amplification dynamics, not voter persuasion. The June 2 result, discussed next, converts this hedge into a finding.

## After the primary: the June 2 result as evidence

The primary was held June 2, 2026. Final results: Bass 34.32% (275,992 votes), Raman 28.55% (229,576), Pratt 25.83% (207,757). Bass and Raman advance to the November runoff; Pratt is eliminated. Pratt held second on election night before late-counted mail ballots broke for Raman “the standard California “blue shift” under universal vote-by-mail, not evidence of irregularity, though such election-night-to-final reversals will themselves become contested material in future cases of this kind.

The result is the strongest available empirical support for the paper’s central hedge. The most media-saturated candidate in the race “3.9-million-view artifacts, elite amplification, five simultaneous national press frames” finished third behind two conventional candidates, one of whom matched his fundraising without any AI ecosystem at all. If supporter-generated AI media converted reach into votes at any substantial rate, this is not the result the configuration would have produced. Attention parity without outcome parity is no longer a hypothesis; it is the observed pattern.

And yet the result complicates as much as it confirms, in one specific way. Pratt’s 25.83% slightly overperformed his final May polling, and *The Hill* assessed it as “by a wide margin the high-water mark for a GOP-aligned candidate since L.A. moved to its current system 14 years ago.” A registered Republican in an overwhelmingly Democratic city drew over 200,000 primary votes. The persuasion-versus-consolidation binary that structured the paper’s original hedge cannot comfortably absorb this: a fourteen-year partisan ceiling was not broken by consolidating an existing Republican base, nor is there evidence of mass persuasion of Democratic voters. The result suggests a third category “mobilization. Supporter-generated AI media may neither persuade marginal voters nor merely consolidate committed ones, but instead activate the politically disengaged and the outsider-curious to participate at all: people for whom the artifacts, rather than the platform, were the point of entry into the election. The Pratt case cannot demonstrate this mechanism “turnout decomposition and voter-file analysis would be required” but it can name it, and the persuasion/consolidation/mobilization triad is a more adequate

dependent-variable structure for the 2028 research program than the binary the field inherited from earlier social-media-effects debates.

## Three propositions for 2028

*First*, the production-layer change is irreversible at the technology level, and the Tier II supplier population will grow accordingly “plausibly from a handful of producers per campaign to dozens or hundreds. This is a claim about capability diffusion, not uniform adoption; celebrity candidates with dense parasocial audiences may continue to see far thicker supporter ecosystems than conventional candidates.

*Second*, curation will become a campaign capability rather than a discretionary task. Campaigns receiving large volumes of supporter-generated artifacts will need to triage “recognizing Curran-quality work, flagging cultural-fit failures and labor-frame exposure *before* amplification, and managing the legal status of each repost. Campaigns that treat the Tier II layer as exogenous noise will be out-produced; campaigns that treat it as an extension of paid media will inherit its coordination-law exposure.

*Third*, the attributable-but-deniable structure will require normative or regulatory address. The space scales with supporter volume; platforms, campaigns, regulators, courts, and the press each hold plausible interventions, and as of this writing none has made one. The first enforcement action or advisory opinion on candidate amplification of supporter-generated AI content will be a field-defining event, and scholars should be positioned to study it.

## Conclusion

The 2026 Los Angeles mayoral campaign of Spencer Pratt is a generative case for a configuration of political media production the field will examine in many more cases across the 2026 and 2028 cycles. The case is best understood through Chadwick’s hybrid media system framework, extended with an explicit production layer that recognizes generative AI as an infrastructural change in who can produce broadcast-quality political video “continuous with the citizen-marketing tradition Penney theorized, discontinuous in the production quality it places in citizen hands. The revised typology “two production tiers coupled to an amplification stack “keeps production provenance and circulation role analytically distinct while modeling their interdependence. The configuration’s constraints are now better specified than its powers: two named failure modes, an unresolved attribution-and-accountability problem with a live regulatory dimension, and a demonstrated ceiling “attention parity without outcome parity. What the configuration may genuinely do, and what the field should now test, is mobilize: the fourteen-year partisan high-water mark Pratt’s losing campaign set suggests that its distinctive electoral effect, if it has one, operates on participation rather than preference. We do not yet know whether the configuration on net improves or degrades the quality of political communication. We know that it is here, that it is likely to grow, and that we owe the field better empirical and conceptual purchase on it.



### Cite this work

Rossi, L. (2026). Supporter as studio: Distributed AI media production in the 2026 Los Angeles mayoral race. Conference paper (working draft), NCA 112th Annual Convention, Political Communication Division. <https://lancerossiconsulting.github.io/pratt-study/conference-paper/>