

What the Spencer Pratt campaign means for your 2028 digital strategy

Lance Rossi

The short version

The most-watched political ad in the 2026 Los Angeles mayoral race “a Batman-themed AI video that drew 3.9 million views on X within a day, was amplified by Jeb Bush, Ted Cruz, and Matt Gaetz, and was covered by every major political outlet” was not made by the Pratt campaign. It was made by one Los Angeles AI filmmaker (Charlie Curran, Menace Studios) on a weekend, for free, with no formal campaign relationship. The campaign reposted it. That is the configuration you need to prepare for. This brief is how to do it.

Three things to do this year

1. Stand up a supporter-relations function.

The single highest-leverage organizational move for 2028 is creating an explicit *supporter-content triage* function inside your digital or communications team. Not a side responsibility for the social manager. A named role, with criteria, with a process. Its job is to (a) monitor supporter-generated AI content as it appears, (b) assess each piece against your campaign’s values, message discipline, and cultural-fit standards, and (c) make the repost / don’t-repost decision before the candidate’s account does. The Pratt campaign appears to have made some of these decisions well (the Curran Batman ad) and some poorly (the Latinos Por Pratt “Spencer, Saca La Bassura” video, which the *LA Times*’s Gustavo Arellano correctly identified as a cultural mismatch in a city whose Latino population is overwhelmingly Mexican). The asymmetry between a good repost and a bad one is enormous; the discipline that distinguishes them is a learnable, hireable capability.

2. Recruit your AI-fluent supporters early, but do not commission them.

The legal and reputational benefits of supporter content being *genuinely* supporter-generated rather than secretly campaign-funded are real, and they matter. Identify the AI-fluent creators in your supporter base “they are findable through the obvious channels (YouTube, TikTok, X)” and build a relationship that consists of (a) reposting their work when it’s good, (b) responding when they tag you, (c) giving them access to candidate appearances at the press-friendly tier, and (d) leaving them alone to create what they want to create. Do not pay them. Do not script them. Do not coordinate the content. The configuration’s value to you depends on the supporters being supporters in fact.

3. Build a cultural-fit review layer inside your triage function.

This is the failure mode that will produce the worst bad-news cycles in 2028. Supporter-generated AI content can be cinematically excellent and culturally tone-deaf at the same time. Build into your triage process at least one reviewer per major cultural community in your coalition (language, region, generation) whose explicit job is to flag the cultural-fit failures before you amplify them. The cost of one Arellano-style column on your most-shared Spanish-language ad is much higher than the cost of one part-time reviewer. This is the cheapest insurance you can buy.

Three things to be honest about

Risk 1: authenticity and attribution.

When the candidate reposts a supporter-generated ad, the campaign is not the legal author but is, in practical political terms, an amplifier of authorship. The “fan-made” frame works for now; it will not work indefinitely. Plan for a 2028 cycle in which press, platforms, opponents, and regulators are paying more attention to the campaign-versus-supporter attribution question. Build your repost criteria as if those criteria will be public.

Risk 2: misinformation and violence in supporter imagery.

Karen Bass on CNN called the AI artifacts a “very dangerous trend” and cited specifically a clip showing her and Newsom drowning in a reservoir. Imagery that crosses your campaign’s threshold for political violence “even if a supporter made it” becomes your problem the moment you amplify it. Decide your line in advance. Document it. Make repost decisions against it. Be prepared to publicly decline to amplify supporter content that crosses it. The supporters who matter to you will respect the line. The ones who don’t were not going to be useful collaborators anyway.

Risk 3: the configuration is not a substitute for the ground game.

The supporter-generated AI ad ecosystem is an extraordinary force multiplier on the air war. It does not knock doors. It does not register voters. It does not turn out low-propensity voters in a primary. Karen Bass retains the institutional money and ground-game advantages; the Pratt campaign’s polling improvement is real but coexists with all the conventional disadvantages of an outsider candidacy. Treat the configuration as additive to a serious campaign, not as a replacement for one.

Three things to integrate immediately

- **Integrate triage with paid media planning.** When a supporter ad performs well organically, your paid team should know within hours whether to put dollars behind boosting it or producing a campaign-original on the same theme. The decision protocol matters more than the speed.

- **Integrate triage with surrogate operations.** When the candidate is asked about a supporter-generated ad in a debate or interview, the candidate needs a one-sentence answer ready that distinguishes “I reposted this because I think it’s clever” from “I commissioned this.” The communications shop needs to have decided those distinctions in advance.
- **Integrate triage with field.** The field team is the campaign’s cultural-fit antenna. Field staff in a community see the supporter content circulating in that community before the digital team does. Build a two-way channel so field flags concerning artifacts and digital responds with restraint or amplification accordingly.

Bottom line

In 2028, the campaigns that treat supporter-generated AI media as exogenous noise will be out-produced by the campaigns that build curatorial discipline around it. The campaigns that treat it as a paid-media substitute will be out-priced by the campaigns that integrate it into a serious ground-game-and-air-war strategy. The campaigns that treat it as a curatorial problem “who develop the triage function, the cultural-fit review, the attribution discipline, and the relationship-not-commission model with their AI-fluent supporters” will, on the present evidence, have the best of the configuration’s upside while managing the downside.

The configuration is here. The question for your 2028 strategy is whether you build the capability to use it well or get used by it badly. There is no third option.

This brief synthesizes the empirical material in the accompanying [NCA conference paper](#) and [journal article](#). For methodological notes, sources, and the broader analysis, see those documents. Comments and pushback welcomed.